

sml-ass-naive-bayes-ipynb

April 6, 2023

0.1 Assignment 8: Naive Bayes Classifier

Supervised Machine Learning Lab ### Kaustubh Raykar PRN : 21070126048 AIML A3

Build a Naïve Bayes Classifier to Predict whether income exceeds \$50K/yr based on census data. Also known as “Census Income” dataset. Dataset is attached

0.1.1 Importing Libraries

```
[150]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
```

0.1.2 Upload Dataset

```
[151]: col=['Age', 'Workclass', 'Fnlwgt', 'Education', 'Education_num', 'Marital_status', 'Occupation', 'Relationship', 'Race', 'Sex', 'Marital_status', 'Occupation', 'Relationship', 'Race', 'Sex']

df = pd.read_csv('/content/Naive bayes dataset.csv', names=col)

df.head(6)
```

```
[151]:
```

	Age	Workclass	Fnlwgt	Education	Education_num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	
5	37	Private	284582	Masters	14	

	Marital_status	Occupation	Relationship	Race	Sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

5	Married-civ-spouse	Exec-managerial		Wife	White	Female
	Capital_gain	Capital_loss	Hours_per_week	Native_country	Income	
0	2174	0	40	United-States	<=50K	
1	0	0	13	United-States	<=50K	
2	0	0	40	United-States	<=50K	
3	0	0	40	United-States	<=50K	
4	0	0	40	Cuba	<=50K	
5	0	0	40	United-States	<=50K	

0.1.3 Data Description

```
[152]: print(df.columns)
```

```
Index(['Age', 'Workclass', 'Fnlwgt', 'Education', 'Education_num',
      'Marital_status', 'Occupation', 'Relationship', 'Race', 'Sex',
      'Capital_gain', 'Capital_loss', 'Hours_per_week', 'Native_country',
      'Income'],
      dtype='object')
```

```
[153]: df.dtypes
```

```
[153]: Age                int64
Workclass              object
Fnlwgt                int64
Education              object
Education_num          int64
Marital_status         object
Occupation             object
Relationship           object
Race                  object
Sex                   object
Capital_gain           int64
Capital_loss           int64
Hours_per_week         int64
Native_country         object
Income                 object
dtype: object
```

```
[154]: df = df.drop(['Fnlwgt'], axis=1)
```

```
[155]: df.drop(['Capital_gain', 'Capital_loss', 'Hours_per_week'], axis=1,
              inplace=True)
```

```
[156]: df
```

```
[156]:
```

	Age	Workclass	Education	Education_num	\
0	39	State-gov	Bachelors	13	
1	50	Self-emp-not-inc	Bachelors	13	
2	38	Private	HS-grad	9	
3	53	Private	11th	7	
4	28	Private	Bachelors	13	
...	
32556	27	Private	Assoc-acdm	12	
32557	40	Private	HS-grad	9	
32558	58	Private	HS-grad	9	
32559	22	Private	HS-grad	9	
32560	52	Self-emp-inc	HS-grad	9	

	Marital_status	Occupation	Relationship	Race	\
0	Never-married	Adm-clerical	Not-in-family	White	
1	Married-civ-spouse	Exec-managerial	Husband	White	
2	Divorced	Handlers-cleaners	Not-in-family	White	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	
4	Married-civ-spouse	Prof-specialty	Wife	Black	
...	
32556	Married-civ-spouse	Tech-support	Wife	White	
32557	Married-civ-spouse	Machine-op-inspct	Husband	White	
32558	Widowed	Adm-clerical	Unmarried	White	
32559	Never-married	Adm-clerical	Own-child	White	
32560	Married-civ-spouse	Exec-managerial	Wife	White	

	Sex	Native_country	Income
0	Male	United-States	<=50K
1	Male	United-States	<=50K
2	Male	United-States	<=50K
3	Male	United-States	<=50K
4	Female	Cuba	<=50K
...
32556	Female	United-States	<=50K
32557	Male	United-States	>50K
32558	Female	United-States	<=50K
32559	Male	United-States	<=50K
32560	Female	United-States	>50K

[32561 rows x 11 columns]

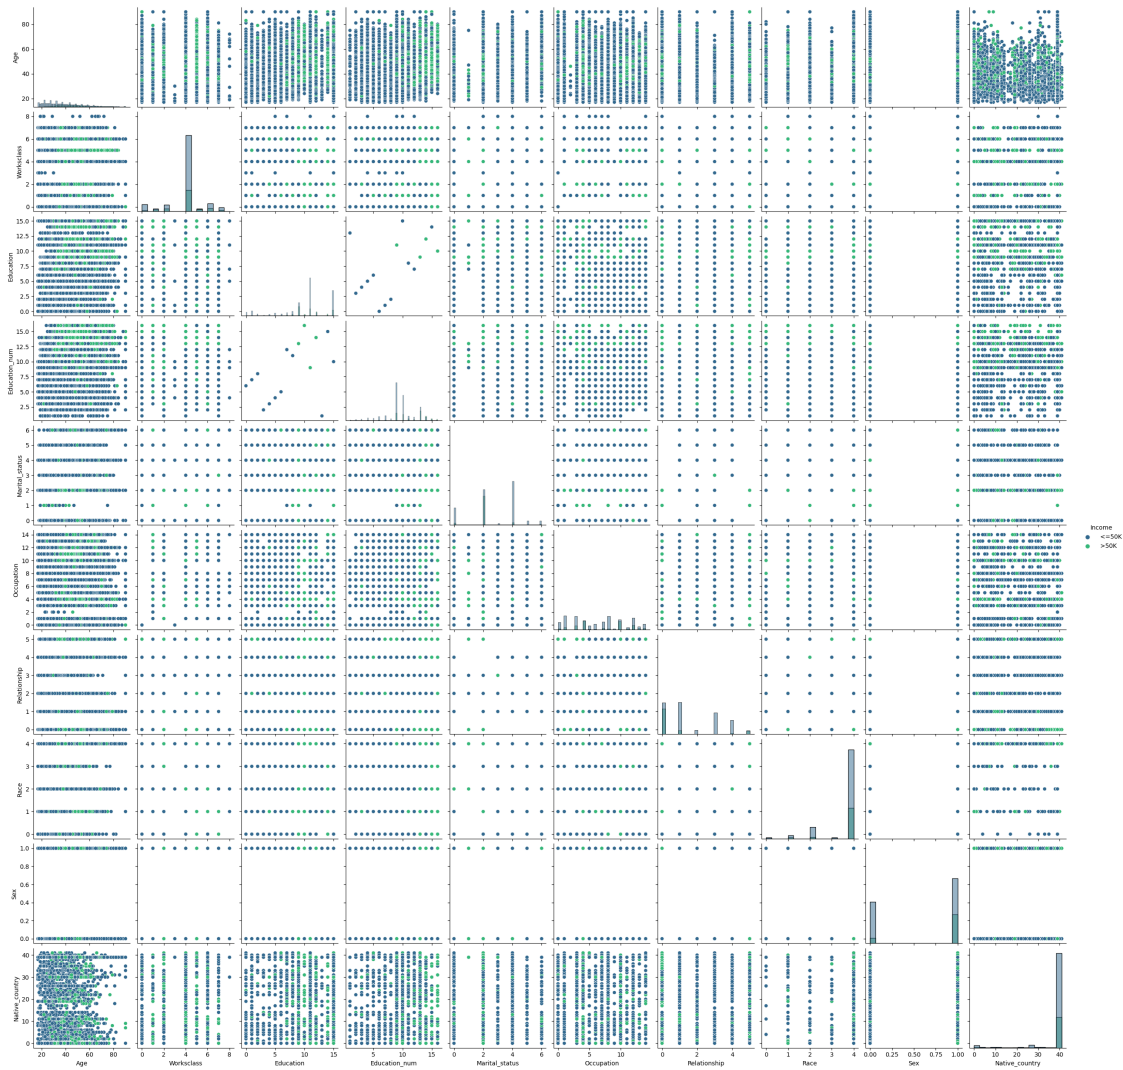
0.1.4 Data Visualisation , EDA

```
[177]: import seaborn as sns

cols = ['Age', 'Workclass', 'Education', 'Education_num', 'Marital_status', 'Occupation', 'Relationship', 'Race', 'Sex', 'Native_country', 'Income']
```


```
sns.pairplot(df[cols], hue='Income', diag_kind='hist', palette='viridis')
```

[177]: <seaborn.axisgrid.PairGrid at 0x7fe464282fa0>



```
[178]: sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
```

[178]: <Axes: >



```
[179]: import plotly.express as px
```

```
fig = px.scatter(df, x="Age", y="Income", color="Sex")
fig.show()
```

```
[181]: fig = px.histogram(df, x="Age", color="Income")
fig.show()
```

```
[182]: fig = px.violin(df, x="Income", y="Education_num", color="Income")
fig.show()
```

```
[183]: print(df['Workclass'].unique())
print(df['Education'].unique())
print(df['Marital_status'].unique())
print(df['Occupation'].unique())
print(df['Relationship'].unique())
print(df['Race'].unique())
print(df['Sex'].unique())
```

```
print(df['Native_country'].unique())
```

```
[7 6 4 1 2 0 5 8 3]
[ 9 11  1 12  6 15  7  8  5 10 14  4  0  3 13  2]
[4 2 0 3 5 1 6]
[ 1  4  6 10  8 12  3 14  5  7 13  0 11  2  9]
[1 0 5 3 4 2]
[4 2 1 0 3]
[1 0]
[39  5 23 19  0 26 35 33 16  9  2 11 20 30 22 31  4  1 37  7 25 36 14 32
  6  8 10 13  3 24 41 29 28 34 38 12 27 40 17 21 18 15]
```

0.1.5 Encoding columns

```
[184]: # Select the categorical columns to be encoded
cat_cols = ['Workclass', 'Education', 'Marital_status', 'Occupation',
            'Relationship', 'Race', 'Sex', 'Native_country']
# Create an instance of LabelEncoder for each column and fit it to the data
label_encoders = {}
for col in cat_cols:
    label_encoders[col] = LabelEncoder()
    df[col] = label_encoders[col].fit_transform(df[col])

# Print the first 5 rows of the encoded data
print(df.head())
```

	Age	Workclass	Education	Education_num	Marital_status	Occupation \
0	39	7	9	13	4	1
1	50	6	9	13	2	4
2	38	4	11	9	0	6
3	53	4	1	7	2	6
4	28	4	9	13	2	10

	Relationship	Race	Sex	Native_country	Income
0	1	4	1	39	<=50K
1	0	4	1	39	<=50K
2	1	4	1	39	<=50K
3	0	2	1	39	<=50K
4	5	2	0	5	<=50K

```
[185]: for col in cat_cols:
        le = label_encoders[col]
        print(f"Column: {col}")
        for i, class_label in enumerate(le.classes_):
            print(f"Label {i}: {class_label}")
        print("\n")
```

Column: Workclass

Label 0: 0
Label 1: 1
Label 2: 2
Label 3: 3
Label 4: 4
Label 5: 5
Label 6: 6
Label 7: 7
Label 8: 8

Column: Education

Label 0: 0
Label 1: 1
Label 2: 2
Label 3: 3
Label 4: 4
Label 5: 5
Label 6: 6
Label 7: 7
Label 8: 8
Label 9: 9
Label 10: 10
Label 11: 11
Label 12: 12
Label 13: 13
Label 14: 14
Label 15: 15

Column: Marital_status

Label 0: 0
Label 1: 1
Label 2: 2
Label 3: 3
Label 4: 4
Label 5: 5
Label 6: 6

Column: Occupation

Label 0: 0
Label 1: 1
Label 2: 2
Label 3: 3
Label 4: 4
Label 5: 5
Label 6: 6

Label 7: 7
Label 8: 8
Label 9: 9
Label 10: 10
Label 11: 11
Label 12: 12
Label 13: 13
Label 14: 14

Column: Relationship

Label 0: 0
Label 1: 1
Label 2: 2
Label 3: 3
Label 4: 4
Label 5: 5

Column: Race

Label 0: 0
Label 1: 1
Label 2: 2
Label 3: 3
Label 4: 4

Column: Sex

Label 0: 0
Label 1: 1

Column: Native_country

Label 0: 0
Label 1: 1
Label 2: 2
Label 3: 3
Label 4: 4
Label 5: 5
Label 6: 6
Label 7: 7
Label 8: 8
Label 9: 9
Label 10: 10
Label 11: 11
Label 12: 12
Label 13: 13
Label 14: 14

Label 15: 15
Label 16: 16
Label 17: 17
Label 18: 18
Label 19: 19
Label 20: 20
Label 21: 21
Label 22: 22
Label 23: 23
Label 24: 24
Label 25: 25
Label 26: 26
Label 27: 27
Label 28: 28
Label 29: 29
Label 30: 30
Label 31: 31
Label 32: 32
Label 33: 33
Label 34: 34
Label 35: 35
Label 36: 36
Label 37: 37
Label 38: 38
Label 39: 39
Label 40: 40
Label 41: 41

```
[186]: # count the number of null values in each column
null_counts = df.isnull().sum()

# print the null counts
print(null_counts)
```

```
Age          0
Workclass    0
Education    0
Education_num 0
Marital_status 0
Occupation   0
Relationship 0
Race         0
Sex          0
Native_country 0
Income       0
dtype: int64
```

```
[187]: df
```

```
[187]:
```

	Age	Workclass	Education	Education_num	Marital_status	Occupation	\
0	39	7	9	13	4	1	
1	50	6	9	13	2	4	
2	38	4	11	9	0	6	
3	53	4	1	7	2	6	
4	28	4	9	13	2	10	
...	
32556	27	4	7	12	2	13	
32557	40	4	11	9	2	7	
32558	58	4	11	9	6	1	
32559	22	4	11	9	4	1	
32560	52	5	11	9	2	4	

	Relationship	Race	Sex	Native_country	Income
0	1	4	1	39	<=50K
1	0	4	1	39	<=50K
2	1	4	1	39	<=50K
3	0	2	1	39	<=50K
4	5	2	0	5	<=50K
...
32556	5	4	0	39	<=50K
32557	0	4	1	39	>50K
32558	4	4	0	39	<=50K
32559	3	4	1	39	<=50K
32560	5	4	0	39	>50K

[32561 rows x 11 columns]

0.1.6 Splitting into x and y

```
[188]: x = df.drop('Income', axis=1)
y = df['Income']
```

```
[189]: x
```

```
[189]:
```

	Age	Workclass	Education	Education_num	Marital_status	Occupation	\
0	39	7	9	13	4	1	
1	50	6	9	13	2	4	
2	38	4	11	9	0	6	
3	53	4	1	7	2	6	
4	28	4	9	13	2	10	
...	
32556	27	4	7	12	2	13	
32557	40	4	11	9	2	7	
32558	58	4	11	9	6	1	

32559	22	4	11	9	4	1
32560	52	5	11	9	2	4

	Relationship	Race	Sex	Native_country
0	1	4	1	39
1	0	4	1	39
2	1	4	1	39
3	0	2	1	39
4	5	2	0	5
...
32556	5	4	0	39
32557	0	4	1	39
32558	4	4	0	39
32559	3	4	1	39
32560	5	4	0	39

[32561 rows x 10 columns]

[190]: y

```
[190]: 0      <=50K
      1      <=50K
      2      <=50K
      3      <=50K
      4      <=50K
      ...
      32556    <=50K
      32557    >50K
      32558    <=50K
      32559    <=50K
      32560    >50K
      Name: Income, Length: 32561, dtype: object
```

0.1.7 Splitting the data into training and testing sets

```
[191]: # Split data into train and test sets
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳ random_state=42)
```

```
[192]: # Print the size of the train and test sets
print('Size of x_train:', x_train.shape)
print('Size of x_test:', x_test.shape)
print('Size of y_train:', y_train.shape)
print('Size of y_test:', y_test.shape)
```

```
Size of x_train: (26048, 10)
Size of x_test: (6513, 10)
```

Size of y_train: (26048,)
Size of y_test: (6513,)

0.1.8 Import and apply Naive Bayes model

```
[193]: from sklearn.naive_bayes import GaussianNB
classifier= GaussianNB()
classifier.fit(x_train, y_train)
```

```
[193]: GaussianNB()
```

```
[194]: y_pred = classifier.predict(x_test)
```

0.1.9 Classification Results of our Model

```
[195]: # Making the Confusion Matrix
from sklearn.metrics import classification_report, confusion_matrix

# Compute the confusion matrix and classification report
cm = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

# Print the results
print("Confusion Matrix:\n", cm)
print("\nClassification Report:\n", report)
```

Confusion Matrix:

```
[[3988  954]
 [ 469 1102]]
```

Classification Report:

	precision	recall	f1-score	support
<=50K	0.89	0.81	0.85	4942
>50K	0.54	0.70	0.61	1571
accuracy			0.78	6513
macro avg	0.72	0.75	0.73	6513
weighted avg	0.81	0.78	0.79	6513

```
[196]: # Format the output of ac and cm
output = 'The accuracy is {:.2f}%\n\nThe confusion matrix is:\n{'
output = output.format(ac*100, cm)

# Print the output
print(output)
```

The accuracy is 78.15%

The confusion matrix is:

```
[[3988  954]
 [ 469 1102]]
```

0.1.10 Confusion Matrix graph


```
[197]: import seaborn as sns

# Calculate the confusion matrix and accuracy score
ac = accuracy_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)

# Plot the confusion matrix using seaborn
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')

# Add labels and title to the plot
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix (Accuracy={:.2f}%)'.format(ac*100))

# Show the plot
plt.show()
```



0.1.11 Thank you

[197]: